

Problem4

April 19, 2022

1 Problem 4 - Optimus, Learning and Resource models, Performance-cost tradeoffs(30 points)

Peng et al. proposed Optimus scheduler for deep learning clusters which makes use of a predictive model to estimate the remaining time of a training job. Optimus assumes a parameter-server architecture for distributed training where synchronization between parameter server(s) and workers happen after every training step. The time taken to complete one training step on a worker includes the time for doing forward propagation (i.e., loss computation) and backward propagation (i.e., gradients computation) at the worker, the worker pushing gradients to parameter servers, parameter servers updating parameters, and the worker pulling updated parameters from parameter servers, plus extra communication overhead.

The predictive model proposed in Optimus is based on two sub-models, one to model the training loss as a function of number of steps and the other to model the training speed (training steps per unit time) as a function of resources (number of workers and parameter servers). The training loss model is given by Equation (1) in the paper. It has three parameters θ_0 , θ_1 , and θ_2 that needs to be estimated from the data.

I collaborate with:

Name	Contribution
Aneri Patel	Resnet-18, 20
Kunal Kulkarni	Resnet-32
Shriya Jain	Resnet-44
Xiang Pang	Resnet-56

```
[ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.nn.init as init
import torchvision
from torch.autograd import Variable
import torchvision.transforms as transforms
import torch.optim as optim

networks = ['resnet18', 'resnet20', 'resnet32', 'resnet44', 'resnet56', 'resnet50']
```

```

GPU = torch.cuda.get_device_name(0)
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
# print(device, "\n\n")

BATCH_SIZE = 128

train_transforms = transforms.Compose(
    [transforms.ToTensor(), transforms.RandomHorizontalFlip(), torchvision.
      transforms.RandomRotation((-15, +15)),
      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

train_dataset = torchvision.datasets.CIFAR10(root='./data', train=True,
    ↳download=True, transform = train_transforms)
test_dataset = torchvision.datasets.CIFAR10(root='./data', train=False,
    ↳download=True, transform=test_transforms)

trainloader = torch.utils.data.DataLoader(train_dataset,
    ↳batch_size=BATCH_SIZE, shuffle=True, num_workers=2)
testloader = torch.utils.data.
    ↳DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False, num_workers=2)

classes = test_dataset.classes
# print(classes)

import torch
import torch.nn as nn
import torch.nn.functional as F

class Block_Plain(nn.Module):
    def __init__(self, num_layers, in_channels, out_channels,
    ↳identity_downsample=None, stride=1):
        assert num_layers in [18, 34, 50, 101, 152], "should be a a valid
    ↳architecture"
        super(Block_Plain, self).__init__()
        self.num_layers = num_layers
        if self.num_layers > 34:
            self.expansion = 4
        else:
            self.expansion = 1
        # ResNet50, 101, and 152 include additional layer of 1x1 kernels

```

```

        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=1,
↪stride=1, padding=0)
        self.bn1 = nn.BatchNorm2d(out_channels)
        if self.num_layers > 34:
            self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3,
↪stride=stride, padding=1)
        else:
            # for ResNet18 and 34, connect input directly to (3x3) kernel (skip
↪first (1x1))
            self.conv2 = nn.Conv2d(in_channels, out_channels, kernel_size=3,
↪stride=stride, padding=1)
            self.bn2 = nn.BatchNorm2d(out_channels)
            self.conv3 = nn.Conv2d(out_channels, out_channels * self.expansion,
↪kernel_size=1, stride=1, padding=0)
            self.bn3 = nn.BatchNorm2d(out_channels * self.expansion)
            self.relu = nn.ReLU()
            self.identity_downsample = identity_downsample

    def forward(self, x):
        identity = x
        if self.num_layers > 34:
            x = self.conv1(x)
            x = self.bn1(x)
            x = self.relu(x)
        x = self.conv2(x)
        x = self.bn2(x)
        x = self.relu(x)
        x = self.conv3(x)
        x = self.bn3(x)

        if self.identity_downsample is not None:
            identity = self.identity_downsample(identity)

        x += identity
        x = self.relu(x)
        return x

class ResNet_Plain(nn.Module):
    def __init__(self, num_layers, block, image_channels, num_classes):
        assert num_layers in [18, 34, 50, 101, 152], f'ResNet{num_layers}:
↪Unknown architecture! Number of layers has ' \
                                                    f'to be 18, 34, 50, 101,
↪or 152 '
        super(ResNet_Plain, self).__init__()
        # From table 1 in ResNet paper

```

```

    if num_layers < 50:
        self.expansion = 1
    else:
        self.expansion = 4
    if num_layers == 18:
        layers = [2, 2, 2, 2]
    elif num_layers == 34 or num_layers == 50:
        layers = [3, 4, 6, 3]
    elif num_layers == 101:
        layers = [3, 4, 23, 3]
    else:
        layers = [3, 8, 36, 3]

    self.in_channels = 64
    self.conv1 = nn.Conv2d(image_channels, 64, kernel_size=7, stride=2,
↪padding=3)
    self.bn1 = nn.BatchNorm2d(64)
    self.relu = nn.ReLU()
    self.maxpool = nn.MaxPool2d(kernel_size=3, stride=2, padding=1)

    # ResNetLayers
    self.layer1 = self.make_layers(num_layers, block, layers[0],
↪intermediate_channels=64, stride=1)
    self.layer2 = self.make_layers(num_layers, block, layers[1],
↪intermediate_channels=128, stride=2)
    self.layer3 = self.make_layers(num_layers, block, layers[2],
↪intermediate_channels=256, stride=2)
    self.layer4 = self.make_layers(num_layers, block, layers[3],
↪intermediate_channels=512, stride=2)

    self.avgpool = nn.AdaptiveAvgPool2d((1, 1))
    self.fc = nn.Linear(512 * self.expansion, num_classes)

    def forward(self, x):
        x = self.conv1(x)
        x = self.bn1(x)
        x = self.relu(x)
        x = self.maxpool(x)

        x = self.layer1(x)
        x = self.layer2(x)
        x = self.layer3(x)
        x = self.layer4(x)

        x = self.avgpool(x)
        x = x.reshape(x.shape[0], -1)
        x = self.fc(x)

```

```

        return x

    def make_layers(self, num_layers, block, num_residual_blocks,
        ↪intermediate_channels, stride):
        layers = []

        identity_downsample = nn.Sequential(nn.Conv2d(self.in_channels,
        ↪intermediate_channels*self.expansion, kernel_size=1, stride=stride),
            nn.
        ↪BatchNorm2d(intermediate_channels*self.expansion))

        layers.append(block(num_layers, self.in_channels,
        ↪intermediate_channels, identity_downsample, stride))
        self.in_channels = intermediate_channels * self.expansion # 256

        for i in range(num_residual_blocks - 1):

            layers.append(block(num_layers, self.in_channels,
        ↪intermediate_channels)) # 256 -> 64, 64*4 (256) again

        return nn.Sequential(*layers)

class block(nn.Module):
    def __init__(self, filters, subsample=False):
        super().__init__()

        s = 0.5 if subsample else 1.0

        self.conv1 = nn.Conv2d(int(filters*s), filters, kernel_size=3,
            stride=int(1/s), padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(filters, track_running_stats=True)
        self.relu1 = nn.ReLU()
        self.conv2 = nn.Conv2d(filters, filters, kernel_size=3, stride=1,
        ↪padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(filters, track_running_stats=True)
        self.relu2 = nn.ReLU()

        # Shortcut downsampling
        self.downsample = nn.AvgPool2d(kernel_size=1, stride=2)

    def shortcut(self, z, x):

        if x.shape != z.shape:
            d = self.downsample(x)
            p = torch.mul(d, 0)

```

```

        return z + torch.cat((d, p), dim=1)
    else:
        return z + x

def forward(self, x, shortcuts=False):
    z = self.conv1(x)
    z = self.bn1(z)
    z = self.relu1(z)

    z = self.conv2(z)
    z = self.bn2(z)

    if shortcuts:
        z = self.shortcut(z, x)

    z = self.relu2(z)

    return z

class ResNet(nn.Module):
    def __init__(self, n, shortcuts=True):
        super().__init__()
        self.shortcuts = shortcuts

        # Input
        self.convIn = nn.Conv2d(3, 16, kernel_size=3, stride=1, padding=1,
↪ bias=False)
        self.bnIn = nn.BatchNorm2d(16, track_running_stats=True)
        self.relu = nn.ReLU()

        # Stack1
        self.stack1 = nn.ModuleList([block(16, subsample=False) for _ in
↪ range(n)])

        # Stack2
        self.stack2a = block(32, subsample=True)
        self.stack2b = nn.ModuleList([block(32, subsample=False) for _ in
↪ range(n-1)])

        # Stack3
        self.stack3a = block(64, subsample=True)
        self.stack3b = nn.ModuleList([block(64, subsample=False) for _ in
↪ range(n-1)])

        self.avgpool = nn.AdaptiveAvgPool2d((1, 1))

```

```

self.fcOut    = nn.Linear(64, 10, bias=True)
self.softmax = nn.LogSoftmax(dim=-1)

def forward(self, x):
    z = self.convIn(x)
    z = self.bnIn(z)
    z = self.relu(z)

    for l in self.stack1: z = l(z, shortcuts=self.shortcuts)

    z = self.stack2a(z, shortcuts=self.shortcuts)
    for l in self.stack2b:
        z = l(z, shortcuts=self.shortcuts)

    z = self.stack3a(z, shortcuts=self.shortcuts)
    for l in self.stack3b:
        z = l(z, shortcuts=self.shortcuts)

    z = self.avgpool(z)
    z = z.view(z.size(0), -1)
    z = self.fcOut(z)
    return self.softmax(z)

def eval_model(val_loader):
    acc = 0
    running_loss = 0.0
    for i, data in enumerate(val_loader, 0):
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)

        outputs = net(inputs)
        loss = criterion(outputs, labels)
        acc += (torch.sum(labels == outputs.argmax(dim=1))/BATCH_SIZE)
        vf.write('%d,%5d,%7f,%7f\n' % (epoch + 1, i + 1, loss.item(), acc/
↪(i+1)))
        running_loss += loss.item()
    print('Epoch: %d, Step: %5d Val loss: %.9f, Acc: %.9f' %
          (epoch + 1, i + 1, running_loss / 50, acc/(i+1)))
    running_loss = 0.0

img_channels = 3
num_classes = 10

```

```

def create_model(n_layers, num_classes=10):
    if n_layers==18:
        net = ResNet_Plain(18, Block_Plain, img_channels, num_classes)
    elif n_layers ==20:
        net = ResNet(3, shortcuts=True)
    elif n_layers ==32:
        net = ResNet(5, shortcuts=True)
    elif n_layers ==44:
        net = ResNet(7, shortcuts=True)
    elif n_layers ==56:
        net = ResNet(9, shortcuts=True)
    # for question 4.2
    elif n_layers==50:
        net = ResNet_Plain(50, Block_Plain, img_channels, num_classes)
    return net

EPOCH = 350

lr = 0.1 # authors cite 0.1
momentum = 0.9
weight_decay = 0.0001
milestones = [82, 123]
gamma = 0.1

# criterion = torch.nn.NLLLoss()

criterion = nn.CrossEntropyLoss()

layers = [18,20,32,44,56,50]
networks = ["resnet"+str(i) for i in layers]

for NETWORK_ID in range(len(layers)):

    print(networks[NETWORK_ID])

    net = create_model(layers[NETWORK_ID]).to(device)
    optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)

    # optimizer = optim.SGD(net.parameters(), lr=lr, momentum=momentum,
    # ↪weight_decay=weight_decay)
    # scheduler = optim.lr_scheduler.MultiStepLR(optimizer,
    # ↪milestones=milestones, gamma=gamma)

    f = open(networks[NETWORK_ID]+"_"+GPU+".csv", "w")

```



```

f.write('epoch, step, loss, acc\n')

vf = open(networks[NETWORK_ID]+"_"+GPU+"_val.csv", "w")
vf.write('epoch, step, loss, acc\n')

for epoch in range(EPOCH):
    acc = 0
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        acc += (torch.sum(labels == outputs.argmax(dim=1))/BATCH_SIZE)
        f.write('%d,%5d,%7f,%7f\n' % (epoch + 1, i + 1, loss.item(), acc/
↪(i+1)))
        running_loss += loss.item()

    if i % 50 == 0: # print every 50 iterations
        print('Epoch: %d, Step: %5d loss: %.9f, Acc: %.9f' %
              (epoch + 1, i + 1, running_loss / 50, acc/(i+1)))
        running_loss = 0.0
        eval_model(testloader)

#         if epoch % 50 == 49:
#             torch.save(net.
↪state_dict(), networks[NETWORK_ID]+"_"+GPU+"_"+str(epoch+1)+".pt")

f.close()
vf.close()
print('Finished Training '+ networks[NETWORK_ID])

```

1.1 1

The first step is to generate data for predictive model calibration. You will train Resnet models with different number of layers (18, 20, 32, 44, 56) each with 3 different GPU types (K80, RTX8000, V100).

For these runs you will use CIFAR10, a batch size of 128, and run each job for 350 epochs.

1.2 Read Results

```
[2]: import pandas as pd
import numpy as np
import os
```

```
[3]: res_path = "./problem4/results/"
for file in os.listdir(res_path):
    if file.endswith(".csv"):
        df = pd.read_csv(res_path + file)
```

```
[4]: df
```

```
[4]:
```

	epoch	step	loss	acc
0	1	1	2.494066	0.109375
1	1	2	2.732457	0.113281
2	1	3	2.677715	0.093750
3	1	4	2.687866	0.105469
4	1	5	2.569639	0.107813
...
221195	350	75	1.053012	0.761042
221196	350	76	1.522796	0.761308
221197	350	77	1.956427	0.761161
221198	350	78	1.571791	0.760918
221199	350	79	2.602238	0.752275

[221200 rows x 4 columns]

```
[5]: from scipy.optimize import curve_fit
```

```
[6]: def func(k, b0, b1, b2):
    loss = 1 / (b0 * k + b1) + b2
    return loss
```

```
[7]: def get_model_parameters(df):
    df.columns = [s.strip() for s in df.columns]
    x = df["real_step"]
    y = df["loss"]
    popt, pcov = curve_fit(func, x, y, bounds=(0, [np.inf, np.inf, np.inf]))
    # print(popt, pcov)
    return popt, pcov
# popt, pcov = get_model_parameters(df)
# b0, b1, b2 = popt
```

```
[8]: def parse_model_name(name):
    name = name.replace(".csv", "")
    model_name = name.split("_")[0]
```

```

gpu_name = name.split("_")[1]
split = "val" if len(name.split("_")) == 3 else "train"
return model_name, gpu_name, split

```

```

[82]: res_list = []
res_path = "./problem4/results/"
for file in os.listdir(res_path):
    if file.endswith(".csv"):
        df = pd.read_csv(res_path + file)
        df.columns = [s.strip() for s in df.columns]
        # print(df)
        epoch_steps = max(df[df["epoch"]==1]["step"])
        df["real_step"] = (df["epoch"] - 1) * epoch_steps + df["step"]

        # first step achieve 92 accuracy
        t = df[df["acc"] > 0.92]["real_step"]
        e = df[df["acc"] > 0.92]["epoch"]

        if len(t) > 0:
            t = t.values[0]
            e = e.values[0]
            loss_92acc = df[df["real_step"] == t]["loss"].values[0]
        else:
            t = -1
            e = -1
            loss_92acc = -1

        # print(t)
        # print(epoch_steps)
        # df["real_step"] = (df["epoch"] == 1)
        popt, pcov = get_model_parameters(df)
        model_name, gpu_name, split = parse_model_name(file)
        b0, b1, b2 = popt
        res_list.append({"model": model_name, "gpu": gpu_name, "split": split,
        ↪ "b0": b0, "b1": b1, "b2": b2, "epoch_to_92acc": e, "step_to_92acc": t,
        ↪ "loss_92acc": loss_92acc, "epoch_steps": epoch_steps})
res_df = pd.DataFrame.from_records(res_list)

```

```
[83]: res_df
```

```

[83]:      model      gpu  split      b0      b1 \
0  resnet56  NVIDIA A100-PCIE-40GB    val    0.007232  5.082540e-01
1  resnet18   Tesla V100-PCIE-32GB    val  15074.454974  4.842781e-15
2  resnet50    Quadro RTX 8000  train    0.000237  4.595429e-01
3  resnet56    Quadro RTX 8000    val    0.006419  5.410299e-01
4  resnet50   Tesla V100-PCIE-32GB  train    0.000233  4.650595e-01
5  resnet44   Tesla V100-PCIE-32GB    val    0.006194  6.156885e-01
6  resnet50  NVIDIA A100-PCIE-40GB  train    0.000236  4.631299e-01

```

7	resnet18	NVIDIA A100-PCIE-40GB	train	0.000271	5.496289e-01
8	resnet20	NVIDIA A100-PCIE-40GB	train	0.000117	6.774311e-01
9	resnet20	Tesla V100-PCIE-32GB	train	0.000116	6.772954e-01
10	resnet44	Quadro RTX 8000	val	0.005784	5.642377e-01
11	resnet20	Tesla V100-PCIE-32GB	val	0.003309	6.152716e-01
12	resnet56	Quadro RTX 8000	train	0.000181	5.943588e-01
13	resnet44	Quadro RTX 8000	train	0.000164	6.297342e-01
14	resnet32	Tesla V100-PCIE-32GB	train	0.000137	6.529592e-01
15	resnet18	Quadro RTX 8000	val	14190.355527	1.761915e-14
16	resnet50	NVIDIA A100-PCIE-40GB	val	15077.103608	1.228839e-14
17	resnet32	Quadro RTX 8000	train	0.000138	6.542840e-01
18	resnet44	NVIDIA A100-PCIE-40GB	val	0.006523	6.016103e-01
19	resnet32	NVIDIA A100-PCIE-40GB	train	0.000135	6.639689e-01
20	resnet56	Tesla V100-PCIE-32GB	val	0.006986	5.944430e-01
21	resnet20	NVIDIA A100-PCIE-40GB	val	0.002869	6.172305e-01
22	resnet18	Tesla V100-PCIE-32GB	train	0.000276	5.409715e-01
23	resnet18	Quadro RTX 8000	train	0.000274	5.473714e-01
24	resnet32	Tesla V100-PCIE-32GB	val	0.004756	5.696338e-01
25	resnet20	Quadro RTX 8000	val	0.003459	6.028435e-01
26	resnet50	Quadro RTX 8000	val	15963.918072	4.808917e-14
27	resnet32	NVIDIA A100-PCIE-40GB	val	0.004970	5.841725e-01
28	resnet56	NVIDIA A100-PCIE-40GB	train	0.000179	6.156932e-01
29	resnet56	Tesla V100-PCIE-32GB	train	0.000168	6.305453e-01
30	resnet50	Tesla V100-PCIE-32GB	val	29266.171362	7.422585e-07
31	resnet32	Quadro RTX 8000	val	0.005167	5.939095e-01
32	resnet44	Tesla V100-PCIE-32GB	train	0.000163	6.476243e-01
33	resnet44	NVIDIA A100-PCIE-40GB	train	0.000152	6.525351e-01
34	resnet20	Quadro RTX 8000	train	0.000119	6.714442e-01
35	resnet18	NVIDIA A100-PCIE-40GB	val	14406.108425	1.890440e-17

	b2	epoch_to_92acc	step_to_92acc	loss_92acc	epoch_steps
0	5.378531e-01	48	3714	0.310703	79
1	1.489328e+00	-1	-1	-1.000000	79
2	5.166572e-18	35	13295	0.245837	391
3	5.360027e-01	64	4978	0.439552	79
4	6.178933e-18	28	10558	0.258628	391
5	5.165396e-01	60	4662	0.326317	79
6	4.789917e-18	32	12122	0.228549	391
7	1.109550e-20	27	10167	0.227245	391
8	3.629090e-02	58	22288	0.308384	391
9	3.292359e-02	65	25025	0.278525	391
10	5.036776e-01	48	3714	0.314827	79
11	4.710460e-01	98	7664	0.398463	79
12	7.921437e-18	43	16423	0.234778	391
13	3.023622e-17	45	17205	0.227768	391
14	4.390594e-16	40	15251	0.213712	391
15	1.470187e+00	-1	-1	-1.000000	79

16	1.731670e+00	-1	-1	-1.000000	79
17	1.229320e-14	49	18769	0.211670	391
18	5.284625e-01	40	3082	0.341347	79
19	8.933287e-16	63	24243	0.203317	391
20	5.457043e-01	95	7427	0.436796	79
21	4.385978e-01	90	7032	0.433734	79
22	5.573940e-20	29	10949	0.209562	391
23	2.673371e-20	31	11731	0.171517	391
24	4.845599e-01	87	6795	0.392498	79
25	4.675423e-01	97	7585	0.361966	79
26	1.763151e+00	-1	-1	-1.000000	79
27	4.905964e-01	33	2529	0.410197	79
28	1.386363e-18	40	15250	0.208657	391
29	1.664642e-18	28	10558	0.265990	391
30	1.779010e+00	-1	-1	-1.000000	79
31	5.055491e-01	66	5136	0.368251	79
32	8.308141e-16	43	16423	0.201995	391
33	1.077400e-17	43	16423	0.235092	391
34	3.446392e-02	52	19942	0.288194	391
35	1.435279e+00	-1	-1	-1.000000	79

1.2.1 Select Training Split

```
[84]: train_res_df = res_df[res_df["split"] == "train"]
```

```
[85]: train_res_df
```

```
[85]:
```

	model	gpu	split	b0	b1	b2	\
2	resnet50	Quadro RTX 8000	train	0.000237	0.459543	5.166572e-18	
4	resnet50	Tesla V100-PCIE-32GB	train	0.000233	0.465060	6.178933e-18	
6	resnet50	NVIDIA A100-PCIE-40GB	train	0.000236	0.463130	4.789917e-18	
7	resnet18	NVIDIA A100-PCIE-40GB	train	0.000271	0.549629	1.109550e-20	
8	resnet20	NVIDIA A100-PCIE-40GB	train	0.000117	0.677431	3.629090e-02	
9	resnet20	Tesla V100-PCIE-32GB	train	0.000116	0.677295	3.292359e-02	
12	resnet56	Quadro RTX 8000	train	0.000181	0.594359	7.921437e-18	
13	resnet44	Quadro RTX 8000	train	0.000164	0.629734	3.023622e-17	
14	resnet32	Tesla V100-PCIE-32GB	train	0.000137	0.652959	4.390594e-16	
17	resnet32	Quadro RTX 8000	train	0.000138	0.654284	1.229320e-14	
19	resnet32	NVIDIA A100-PCIE-40GB	train	0.000135	0.663969	8.933287e-16	
22	resnet18	Tesla V100-PCIE-32GB	train	0.000276	0.540971	5.573940e-20	
23	resnet18	Quadro RTX 8000	train	0.000274	0.547371	2.673371e-20	
28	resnet56	NVIDIA A100-PCIE-40GB	train	0.000179	0.615693	1.386363e-18	
29	resnet56	Tesla V100-PCIE-32GB	train	0.000168	0.630545	1.664642e-18	
32	resnet44	Tesla V100-PCIE-32GB	train	0.000163	0.647624	8.308141e-16	
33	resnet44	NVIDIA A100-PCIE-40GB	train	0.000152	0.652535	1.077400e-17	
34	resnet20	Quadro RTX 8000	train	0.000119	0.671444	3.446392e-02	

	epoch_to_92acc	step_to_92acc	loss_92acc	epoch_steps
2	35	13295	0.245837	391
4	28	10558	0.258628	391
6	32	12122	0.228549	391
7	27	10167	0.227245	391
8	58	22288	0.308384	391
9	65	25025	0.278525	391
12	43	16423	0.234778	391
13	45	17205	0.227768	391
14	40	15251	0.213712	391
17	49	18769	0.211670	391
19	63	24243	0.203317	391
22	29	10949	0.209562	391
23	31	11731	0.171517	391
28	40	15250	0.208657	391
29	28	10558	0.265990	391
32	43	16423	0.201995	391
33	43	16423	0.235092	391
34	52	19942	0.288194	391

1.3 2

We next study how the learned parameters, 0, 1, and 2, change with the type of GPUs and the size of network. Use a regression model on the data from 15 models to predict the value of these parameters as a function of number of layers in Resnet and GPU type. From these regression model predict the training loss curve for Resnet-50. Note that we are effectively doing prediction for a predictive model. To verify how good is this prediction, you will train Resnet-50 on a K80, RTX8000, and V100 for target accuracy of 92% and compare the predicted loss curve with the real measurements. Show this comparison in a graph and calculate the percentage error. From the predicted loss curve get the number of epochs needed to achieve 92% accuracy. Observe that there are three curves for three different GPU types, but the number of epochs required to reach a particular accuracy (convergence rate) should be independent of hardware. (8)

```
[13]: gpu_kinds = train_res_df["gpu"].unique()
      # print(gpu_kinds)
      gpu_type_mapping = {t: i for i, t in enumerate(gpu_kinds)}
      # print(gpu_type_mapping)
```

```
[14]: res_df["gpu_type"] = res_df["gpu"].map(gpu_type_mapping)
```

```
[15]: def get_num_layers(model_name):
      num_layers = model_name.replace("resnet", "")
      return int(num_layers)
      res_df["num_layers"] = res_df["model"].map(get_num_layers)
      res_df.to_csv("./problem4/res_df.csv", index=False)
```

```
[16]: res_df = res_df[res_df["split"] == "train"]
```

```
[17]: train_res_df = res_df[res_df["num_layers"] != 50]
      test_res_df = res_df[res_df["num_layers"] == 50]
```

```
[18]: # b0, b1, b2 = reg(num_layers, )
      x0 = train_res_df["num_layers"].values
      x1 = train_res_df["gpu_type"].values
      X = np.array([x0, x1]).T
      y0 = train_res_df["b0"].values
      y1 = train_res_df["b1"].values
      y2 = train_res_df["b2"].values
      Y = np.array([y0, y1, y2]).T
```

```
[33]: X.shape
```

```
[33]: (15, 2)
```

```
[34]: Y.shape
```

```
[34]: (15, 3)
```

```
[36]: import numpy as np
      from sklearn.datasets import load_linnerud
      from sklearn.multioutput import MultiOutputRegressor
      from sklearn.preprocessing import PolynomialFeatures
      from sklearn.linear_model import Ridge
      from sklearn.kernel_ridge import KernelRidge
      from sklearn.pipeline import make_pipeline
      krr = KernelRidge(alpha=1.0)
      krr.fit(X, Y)
      Y_hat = krr.predict(X)
      yy = np.concatenate((Y, Y_hat), axis=1)
      df = pd.DataFrame(yy, columns=["b0", "b1", "b2", "b0_hat", "b1_hat", "b2_hat"])
      df
```

```
[36]:
```

	b0	b1	b2	b0_hat	b1_hat	b2_hat
0	0.000271	0.549629	1.109550e-20	0.000126	0.453489	0.008108
1	0.000117	0.677431	3.629090e-02	0.000133	0.479377	0.008108
2	0.000116	0.677295	3.292359e-02	0.000100	0.369129	0.004058
3	0.000181	0.594359	7.921437e-18	0.000189	0.724866	0.000022
4	0.000164	0.629734	3.023622e-17	0.000148	0.569538	0.000017
5	0.000137	0.652959	4.390594e-16	0.000140	0.524457	0.004063
6	0.000138	0.654284	1.229320e-14	0.000108	0.414209	0.000012
7	0.000135	0.663969	8.933287e-16	0.000173	0.634705	0.008113
8	0.000276	0.540971	5.573940e-20	0.000093	0.343241	0.004057
9	0.000274	0.547371	2.673371e-20	0.000061	0.232993	0.000007
10	0.000179	0.615693	1.386363e-18	0.000254	0.945362	0.008122
11	0.000168	0.630545	1.664642e-18	0.000221	0.835114	0.004072

```

12  0.000163  0.647624  8.308141e-16  0.000181  0.679786  0.004067
13  0.000152  0.652535  1.077400e-17  0.000214  0.790034  0.008118
14  0.000119  0.671444  3.446392e-02  0.000067  0.258881  0.000008

```

```

[37]: # b0, b1, b2 = reg(num_layers, )
x0 = test_res_df["num_layers"].values
x1 = test_res_df["gpu_type"].values
test_X = np.array([x0, x1]).T
y0 = test_res_df["b0"].values
y1 = test_res_df["b1"].values
y2 = test_res_df["b2"].values
test_Y = np.array([y0, y1, y2]).T

```

```

[38]: test_res_df

```

```

[38]:      model      gpu  split      b0      b1      b2  \
2  resnet50  Quadro RTX 8000  train  0.000237  0.459543  5.166572e-18
4  resnet50  Tesla V100-PCIE-32GB  train  0.000233  0.465060  6.178933e-18
6  resnet50  NVIDIA A100-PCIE-40GB  train  0.000236  0.463130  4.789917e-18

      epoch_to_92acc  step_to_92acc  loss_92acc  epoch_steps  gpu_type  \
2                35          13295    0.245837          391          0
4                28          10558    0.258628          391          1
6                32          12122    0.228549          391          2

      num_layers
2              50
4              50
6              50

```

```

[39]: test_X

```

```

[39]: array([[50,  0],
          [50,  1],
          [50,  2]])

```

```

[92]: test_res_df["gpu_type"]

```

```

[92]: 2    0
      4    1
      6    2
      Name: gpu_type, dtype: int64

```

```

[93]: test_Y_hat = krr.predict(test_X)
yy = np.concatenate((test_Y, test_Y_hat), axis=1)
df = pd.DataFrame(yy, columns=["b0", "b1", "b2", "b0_hat", "b1_hat", "b2_hat"])
df["gpu_type"] = test_res_df["gpu_type"].values

```



```

df["num_layers"] = test_res_df["num_layers"].values
df["model"] = test_res_df["model"].values
df["gpu"] = test_res_df["gpu"].values
def get_file_name(model, gpu):
    return model + "_" + gpu + ".csv"
df["file_name"] = get_file_name(df["model"], df["gpu"])
df

```

```

[93]:
      b0      b1      b2  b0_hat  b1_hat  b2_hat  gpu_type  \
0  0.000237  0.459543  5.166572e-18  0.000168  0.647202  0.000019      0
1  0.000233  0.465060  6.178933e-18  0.000201  0.757450  0.004070      1
2  0.000236  0.463130  4.789917e-18  0.000234  0.867698  0.008120      2

      num_layers      model      gpu  \
0           50  resnet50      Quadro RTX 8000
1           50  resnet50  Tesla V100-PCIE-32GB
2           50  resnet50  NVIDIA A100-PCIE-40GB

      file_name
0  resnet50_Quadro RTX 8000.csv
1  resnet50_Tesla V100-PCIE-32GB.csv
2  resnet50_NVIDIA A100-PCIE-40GB.csv

```

```

[94]: df.to_csv("./problem4/2.csv", index=False)

```

1.3.1 Plot the three curves

```

[95]: def func(k, b0, b1, b2):
      loss = 1 / (b0 * k + b1) + b2
      return loss

```

```

[96]: def get_y(x, params):
      return func(x, *params)

```

```

[97]: import matplotlib.pyplot as plt
res_path = "./problem4/results/"
for i in range(3):
    fig = plt.figure()
    temp_df = pd.read_csv(res_path + df["file_name"][i])
    temp_df.columns = [s.strip() for s in temp_df.columns]

    epoch_steps = max(temp_df["step"])
    temp_df["real_step"] = temp_df["epoch"] * epoch_steps + temp_df["step"]
    real_X = temp_df["real_step"].values
    real_Y = temp_df["loss"].values
    plt.plot(real_X, real_Y, label="real")


```

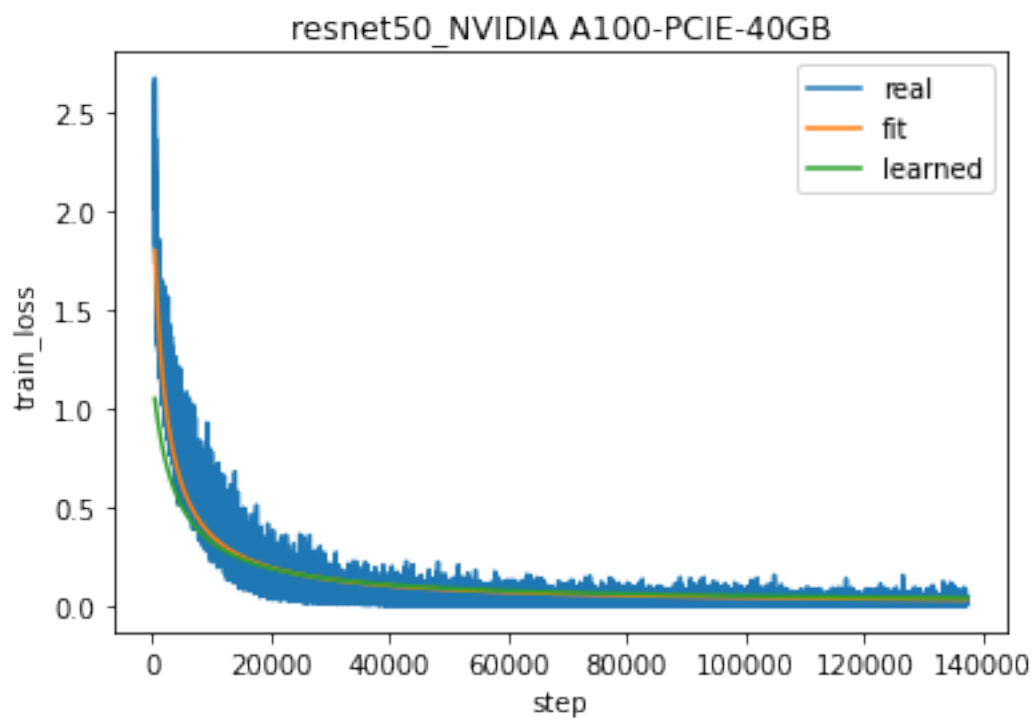
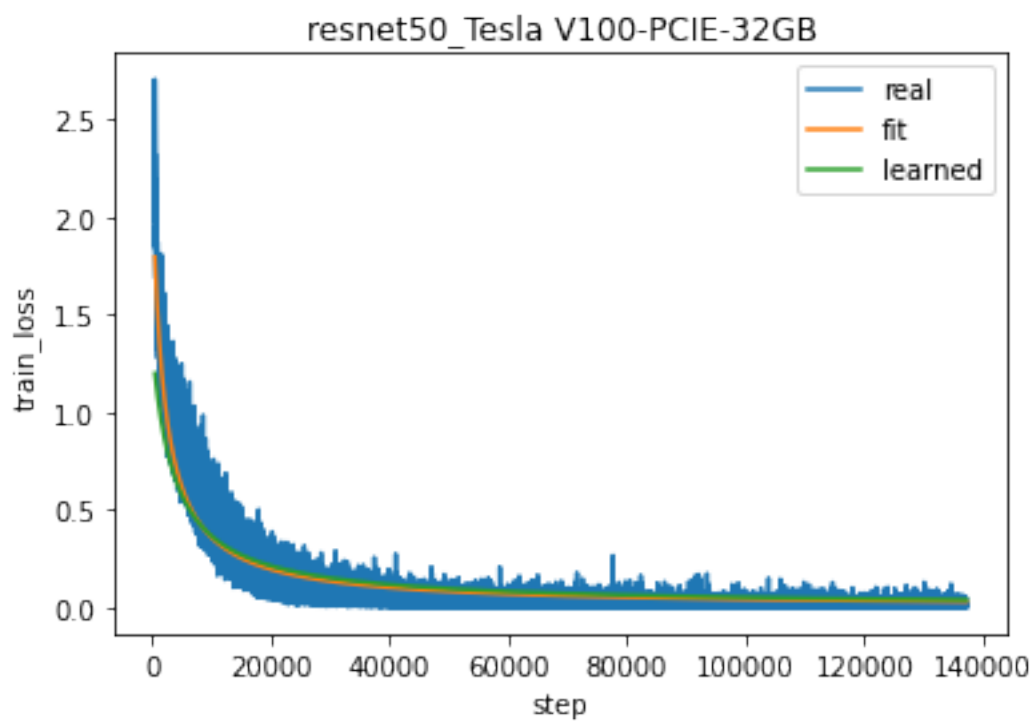
```

b0 = df["b0"][i]
b1 = df["b1"][i]
b2 = df["b2"][i]
fit_X = real_X
fit_Y = get_y(fit_X, (b0, b1, b2))
plt.plot(fit_X, fit_Y, label="fit")

b0_hat = df["b0_hat"][i]
b1_hat = df["b1_hat"][i]
b2_hat = df["b2_hat"][i]
learned_X = real_X
learned_Y = get_y(learned_X, (b0_hat, b1_hat, b2_hat))
plt.plot(learned_X, learned_Y, label="learned")
plt.xlabel("step")
plt.ylabel("train_loss")
plt.title(df["file_name"][i].replace(".csv", ""))
plt.legend()
plt.show()

```





1.3.2 Parameter Error

```
[98]: params_list = ["b0", "b1", "b2"]
for p in params_list:
    df[f"{p}_percentage_error"] = df[f"{p}"] - df[f"{p}_hat"] / df[f"{p}"]
    df[f"{p}_absolut_error"] = df[f"{p}"] - df[f"{p}_hat"]
df
```

```
[98]:
```

	b0	b1	b2	b0_hat	b1_hat	b2_hat	gpu_type	\
0	0.000237	0.459543	5.166572e-18	0.000168	0.647202	0.000019	0	
1	0.000233	0.465060	6.178933e-18	0.000201	0.757450	0.004070	1	
2	0.000236	0.463130	4.789917e-18	0.000234	0.867698	0.008120	2	

	num_layers	model	gpu	\
0	50	resnet50	Quadro RTX 8000	
1	50	resnet50	Tesla V100-PCIE-32GB	
2	50	resnet50	NVIDIA A100-PCIE-40GB	

	file_name	b0_percentage_error	b0_absolut_error	\
0	resnet50_Quadro RTX 8000.csv	-0.710173	0.000069	
1	resnet50_Tesla V100-PCIE-32GB.csv	-0.862069	0.000032	
2	resnet50_NVIDIA A100-PCIE-40GB.csv	-0.992541	0.000002	

	b1_percentage_error	b1_absolut_error	b2_percentage_error	\
0	-0.948817	-0.187659	-3.738277e+12	
1	-1.163657	-0.292390	-6.586309e+14	
2	-1.410422	-0.404568	-1.695219e+15	

	b2_absolut_error
0	-0.000019
1	-0.004070
2	-0.008120

1.3.3 Loss Estimation Percentage Error Over Steps

```
[138]: import matplotlib.pyplot as plt
res_path = "./problem4/results/"
for i in range(3):
    fig = plt.figure()
    temp_df = pd.read_csv(res_path + df["file_name"][i])
    temp_df.columns = [s.strip() for s in temp_df.columns]

    epoch_steps = max(temp_df["step"])
    temp_df["real_step"] = temp_df["epoch"] * epoch_steps + temp_df["step"]
    real_X = temp_df["real_step"].values
    real_Y = temp_df["loss"].values
```



```

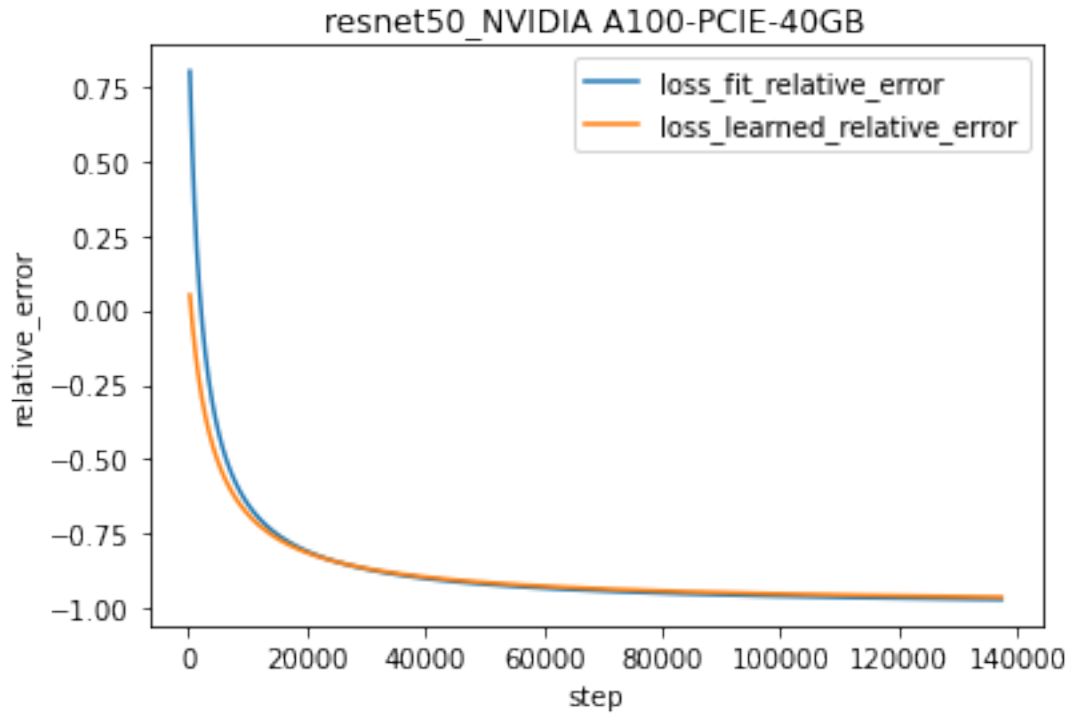
b0 = df["b0"][i]
b1 = df["b1"][i]
b2 = df["b2"][i]
fit_X = real_X
fit_Y = get_y(fit_X, (b0, b1, b2))
fit_percentage_error = fit_Y - real_Y / real_Y
fit_absolut_error = fit_Y - real_Y

b0_hat = df["b0_hat"][i]
b1_hat = df["b1_hat"][i]
b2_hat = df["b2_hat"][i]
learned_X = real_X
learned_Y = get_y(learned_X, (b0_hat, b1_hat, b2_hat))
learned_percentage_error = learned_Y - real_Y / real_Y
learned_absolut_error = learned_Y - real_Y

plt.plot(real_X, fit_percentage_error, label="loss_fit_relative_error")
# plt.plot(real_X, fit_absolut_error, label="real")
plt.plot(real_X, learned_percentage_error,
label="loss_learned_relative_error")
# plt.plot(real_X, learned_absolut_error, label="learned_absolut_error")
# plt.plot
plt.legend()
plt.title(df["file_name"][i].replace(".csv", ""))
plt.xlabel("step")
plt.ylabel("relative_error")
plt.show()

```





1.3.4 Loss Estimation Absolute Error Over Steps

```
[139]: import matplotlib.pyplot as plt
res_path = "./problem4/results/"
for i in range(3):
    fig = plt.figure()
    temp_df = pd.read_csv("./problem4/results/" + df["file_name"][i])
    temp_df.columns = [s.strip() for s in temp_df.columns]

    epoch_steps = max(temp_df["step"])
    temp_df["real_step"] = temp_df["epoch"] * epoch_steps + temp_df["step"]
    real_X = temp_df["real_step"].values
    real_Y = temp_df["loss"].values

    b0 = df["b0"][i]
    b1 = df["b1"][i]
    b2 = df["b2"][i]
    fit_X = real_X
    fit_Y = get_y(fit_X, (b0, b1, b2))
    fit_percentage_error = fit_Y - real_Y / real_Y
    fit_absolut_error = fit_Y - real_Y


    b0_hat = df["b0_hat"][i]
```



```

b1_hat = df["b1_hat"][i]
b2_hat = df["b2_hat"][i]
learned_X = real_X
learned_Y = get_y(learned_X, (b0_hat, b1_hat, b2_hat))
learned_percentage_error = learned_Y - real_Y / real_Y
learned_absolut_error = learned_Y - real_Y

plt.plot(real_X, fit_absolut_error, label="loss_fit_relative_error")
# plt.plot(real_X, fit_absolut_error, label="real")
plt.plot(real_X, learned_absolut_error, label="loss_learned_absolut_error")
# plt.plot(real_X, learned_absolut_error, label="learned_absolut_error")
# plt.plot
plt.legend()
plt.title(df["file_name"][i].replace(".csv", ""))
plt.xlabel("step")
plt.ylabel("absolut_error")
plt.show()

```





1.3.5 Comment

We do not know the clear definition of the percentage error. We plot both the relative error and the absolute error over steps.

For fit, it means we use the real loss fitted curve, and calculate the error. For learned, it means we use the curve parameter learning model to get the curve parameters, and calculate the error.

We can notice that the three types of gpu do not have clear difference in the percentage error.

Epoch to Reach 92% From the predicted loss curve get the number of epochs needed to achieve 92% accuracy.

```
[101]: res_df
```

```
[101]:
```

	model	gpu	split	b0	b1 \
0	resnet56	NVIDIA A100-PCIE-40GB	val	0.007232	5.082540e-01
1	resnet18	Tesla V100-PCIE-32GB	val	15074.454974	4.842781e-15
2	resnet50	Quadro RTX 8000	train	0.000237	4.595429e-01
3	resnet56	Quadro RTX 8000	val	0.006419	5.410299e-01
4	resnet50	Tesla V100-PCIE-32GB	train	0.000233	4.650595e-01
5	resnet44	Tesla V100-PCIE-32GB	val	0.006194	6.156885e-01
6	resnet50	NVIDIA A100-PCIE-40GB	train	0.000236	4.631299e-01
7	resnet18	NVIDIA A100-PCIE-40GB	train	0.000271	5.496289e-01
8	resnet20	NVIDIA A100-PCIE-40GB	train	0.000117	6.774311e-01
9	resnet20	Tesla V100-PCIE-32GB	train	0.000116	6.772954e-01
10	resnet44	Quadro RTX 8000	val	0.005784	5.642377e-01
11	resnet20	Tesla V100-PCIE-32GB	val	0.003309	6.152716e-01
12	resnet56	Quadro RTX 8000	train	0.000181	5.943588e-01
13	resnet44	Quadro RTX 8000	train	0.000164	6.297342e-01
14	resnet32	Tesla V100-PCIE-32GB	train	0.000137	6.529592e-01
15	resnet18	Quadro RTX 8000	val	14190.355527	1.761915e-14
16	resnet50	NVIDIA A100-PCIE-40GB	val	15077.103608	1.228839e-14
17	resnet32	Quadro RTX 8000	train	0.000138	6.542840e-01
18	resnet44	NVIDIA A100-PCIE-40GB	val	0.006523	6.016103e-01
19	resnet32	NVIDIA A100-PCIE-40GB	train	0.000135	6.639689e-01
20	resnet56	Tesla V100-PCIE-32GB	val	0.006986	5.944430e-01
21	resnet20	NVIDIA A100-PCIE-40GB	val	0.002869	6.172305e-01
22	resnet18	Tesla V100-PCIE-32GB	train	0.000276	5.409715e-01
23	resnet18	Quadro RTX 8000	train	0.000274	5.473714e-01
24	resnet32	Tesla V100-PCIE-32GB	val	0.004756	5.696338e-01
25	resnet20	Quadro RTX 8000	val	0.003459	6.028435e-01
26	resnet50	Quadro RTX 8000	val	15963.918072	4.808917e-14
27	resnet32	NVIDIA A100-PCIE-40GB	val	0.004970	5.841725e-01
28	resnet56	NVIDIA A100-PCIE-40GB	train	0.000179	6.156932e-01
29	resnet56	Tesla V100-PCIE-32GB	train	0.000168	6.305453e-01
30	resnet50	Tesla V100-PCIE-32GB	val	29266.171362	7.422585e-07
31	resnet32	Quadro RTX 8000	val	0.005167	5.939095e-01
32	resnet44	Tesla V100-PCIE-32GB	train	0.000163	6.476243e-01

```

33 resnet44  NVIDIA A100-PCIE-40GB  train      0.000152  6.525351e-01
34 resnet20      Quadro RTX 8000  train      0.000119  6.714442e-01
35 resnet18  NVIDIA A100-PCIE-40GB    val  14406.108425  1.890440e-17

```

	b2	epoch_to_92acc	step_to_92acc	loss_92acc	epoch_steps
0	5.378531e-01	48	3714	0.310703	79
1	1.489328e+00	-1	-1	-1.000000	79
2	5.166572e-18	35	13295	0.245837	391
3	5.360027e-01	64	4978	0.439552	79
4	6.178933e-18	28	10558	0.258628	391
5	5.165396e-01	60	4662	0.326317	79
6	4.789917e-18	32	12122	0.228549	391
7	1.109550e-20	27	10167	0.227245	391
8	3.629090e-02	58	22288	0.308384	391
9	3.292359e-02	65	25025	0.278525	391
10	5.036776e-01	48	3714	0.314827	79
11	4.710460e-01	98	7664	0.398463	79
12	7.921437e-18	43	16423	0.234778	391
13	3.023622e-17	45	17205	0.227768	391
14	4.390594e-16	40	15251	0.213712	391
15	1.470187e+00	-1	-1	-1.000000	79
16	1.731670e+00	-1	-1	-1.000000	79
17	1.229320e-14	49	18769	0.211670	391
18	5.284625e-01	40	3082	0.341347	79
19	8.933287e-16	63	24243	0.203317	391
20	5.457043e-01	95	7427	0.436796	79
21	4.385978e-01	90	7032	0.433734	79
22	5.573940e-20	29	10949	0.209562	391
23	2.673371e-20	31	11731	0.171517	391
24	4.845599e-01	87	6795	0.392498	79
25	4.675423e-01	97	7585	0.361966	79
26	1.763151e+00	-1	-1	-1.000000	79
27	4.905964e-01	33	2529	0.410197	79
28	1.386363e-18	40	15250	0.208657	391
29	1.664642e-18	28	10558	0.265990	391
30	1.779010e+00	-1	-1	-1.000000	79
31	5.055491e-01	66	5136	0.368251	79
32	8.308141e-16	43	16423	0.201995	391
33	1.077400e-17	43	16423	0.235092	391
34	3.446392e-02	52	19942	0.288194	391
35	1.435279e+00	-1	-1	-1.000000	79

```
[102]: train_res_df = res_df[res_df["split"] == "train"]
```

```
[103]: train_res_df
```

```
[103]:
```

	model	gpu	split	b0	b1	b2 \
2	resnet50	Quadro RTX 8000	train	0.000237	0.459543	5.166572e-18
4	resnet50	Tesla V100-PCIE-32GB	train	0.000233	0.465060	6.178933e-18
6	resnet50	NVIDIA A100-PCIE-40GB	train	0.000236	0.463130	4.789917e-18
7	resnet18	NVIDIA A100-PCIE-40GB	train	0.000271	0.549629	1.109550e-20
8	resnet20	NVIDIA A100-PCIE-40GB	train	0.000117	0.677431	3.629090e-02
9	resnet20	Tesla V100-PCIE-32GB	train	0.000116	0.677295	3.292359e-02
12	resnet56	Quadro RTX 8000	train	0.000181	0.594359	7.921437e-18
13	resnet44	Quadro RTX 8000	train	0.000164	0.629734	3.023622e-17
14	resnet32	Tesla V100-PCIE-32GB	train	0.000137	0.652959	4.390594e-16
17	resnet32	Quadro RTX 8000	train	0.000138	0.654284	1.229320e-14
19	resnet32	NVIDIA A100-PCIE-40GB	train	0.000135	0.663969	8.933287e-16
22	resnet18	Tesla V100-PCIE-32GB	train	0.000276	0.540971	5.573940e-20
23	resnet18	Quadro RTX 8000	train	0.000274	0.547371	2.673371e-20
28	resnet56	NVIDIA A100-PCIE-40GB	train	0.000179	0.615693	1.386363e-18
29	resnet56	Tesla V100-PCIE-32GB	train	0.000168	0.630545	1.664642e-18
32	resnet44	Tesla V100-PCIE-32GB	train	0.000163	0.647624	8.308141e-16
33	resnet44	NVIDIA A100-PCIE-40GB	train	0.000152	0.652535	1.077400e-17
34	resnet20	Quadro RTX 8000	train	0.000119	0.671444	3.446392e-02

	epoch_to_92acc	step_to_92acc	loss_92acc	epoch_steps
2	35	13295	0.245837	391
4	28	10558	0.258628	391
6	32	12122	0.228549	391
7	27	10167	0.227245	391
8	58	22288	0.308384	391
9	65	25025	0.278525	391
12	43	16423	0.234778	391
13	45	17205	0.227768	391
14	40	15251	0.213712	391
17	49	18769	0.211670	391
19	63	24243	0.203317	391
22	29	10949	0.209562	391
23	31	11731	0.171517	391
28	40	15250	0.208657	391
29	28	10558	0.265990	391
32	43	16423	0.201995	391
33	43	16423	0.235092	391
34	52	19942	0.288194	391

We use the average loss to 92% accuracy as the estimate of loss to reach the accuracy.

```
[59]: threshold = train_res_df["loss_92acc"].mean()
```

```
[104]: threshold
```

```
[104]: 0.23441231666666665
```

```
[110]: df
```

```
[110]:      b0      b1      b2  b0_hat  b1_hat  b2_hat  gpu_type  \
0  0.000237  0.459543  5.166572e-18  0.000168  0.647202  0.000019      0
1  0.000233  0.465060  6.178933e-18  0.000201  0.757450  0.004070      1
2  0.000236  0.463130  4.789917e-18  0.000234  0.867698  0.008120      2

   num_layers  model      gpu  \
0         50  resnet50  Quadro RTX 8000
1         50  resnet50  Tesla V100-PCIE-32GB
2         50  resnet50  NVIDIA A100-PCIE-40GB

      file_name  b0_percentage_error  b0_absolut_error  \
0  resnet50_Quadro RTX 8000.csv      -0.710173      0.000069
1  resnet50_Tesla V100-PCIE-32GB.csv      -0.862069      0.000032
2  resnet50_NVIDIA A100-PCIE-40GB.csv      -0.992541      0.000002

   b1_percentage_error  b1_absolut_error  b2_percentage_error  \
0          -0.948817          -0.187659          -3.738277e+12
1          -1.163657          -0.292390          -6.586309e+14
2          -1.410422          -0.404568          -1.695219e+15

   b2_absolut_error
0          -0.000019
1          -0.004070
2          -0.008120
```

```
[112]: test_res_df
```

```
[112]:      model      gpu  split      b0      b1      b2  \
2  resnet50  Quadro RTX 8000  train  0.000237  0.459543  5.166572e-18
4  resnet50  Tesla V100-PCIE-32GB  train  0.000233  0.465060  6.178933e-18
6  resnet50  NVIDIA A100-PCIE-40GB  train  0.000236  0.463130  4.789917e-18

   epoch_to_92acc  step_to_92acc  loss_92acc  epoch_steps  gpu_type  \
2              35           13295    0.245837           391          0
4              28           10558    0.258628           391          1
6              32           12122    0.228549           391          2

   num_layers
2          50
4          50
6          50
```

```
[130]: df
```

```
[130]:
```

	b0	b1	b2	b0_hat	b1_hat	b2_hat	gpu_type	\
0	0.000237	0.459543	5.166572e-18	0.000168	0.647202	0.000019	0	
1	0.000233	0.465060	6.178933e-18	0.000201	0.757450	0.004070	1	
2	0.000236	0.463130	4.789917e-18	0.000234	0.867698	0.008120	2	

	num_layers	model	gpu	\
0	50	resnet50	Quadro RTX 8000	
1	50	resnet50	Tesla V100-PCIE-32GB	
2	50	resnet50	NVIDIA A100-PCIE-40GB	

	file_name	b0_percentage_error	b0_absolut_error	\
0	resnet50_Quadro RTX 8000.csv	-0.710173	0.000069	
1	resnet50_Tesla V100-PCIE-32GB.csv	-0.862069	0.000032	
2	resnet50_NVIDIA A100-PCIE-40GB.csv	-0.992541	0.000002	

	b1_percentage_error	b1_absolut_error	b2_percentage_error	\
0	-0.948817	-0.187659	-3.738277e+12	
1	-1.163657	-0.292390	-6.586309e+14	
2	-1.410422	-0.404568	-1.695219e+15	

	b2_absolut_error
0	-0.000019
1	-0.004070
2	-0.008120

```
[120]: threshold = test_res_df["loss_92acc"].mean()
```

```
[136]: import matplotlib.pyplot as plt
res_path = "./problem4/results/"
for i in range(3):
    fig = plt.figure()
    temp_df = pd.read_csv("./problem4/results/" + df["file_name"][i])
    temp_df.columns = [s.strip() for s in temp_df.columns]

    epoch_steps = max(temp_df["step"])
    temp_df["real_step"] = temp_df["epoch"] * epoch_steps + temp_df["step"]
    real_X = temp_df["real_step"].values
    real_Y = temp_df["loss"].values
    plt.plot(real_X, real_Y, label="real")

    b0 = df["b0"][i]
    b1 = df["b1"][i]
    b2 = df["b2"][i]
    fit_X = real_X
    fit_Y = get_y(fit_X, (b0, b1, b2))
    plt.plot(fit_X, fit_Y, label="fit")
```

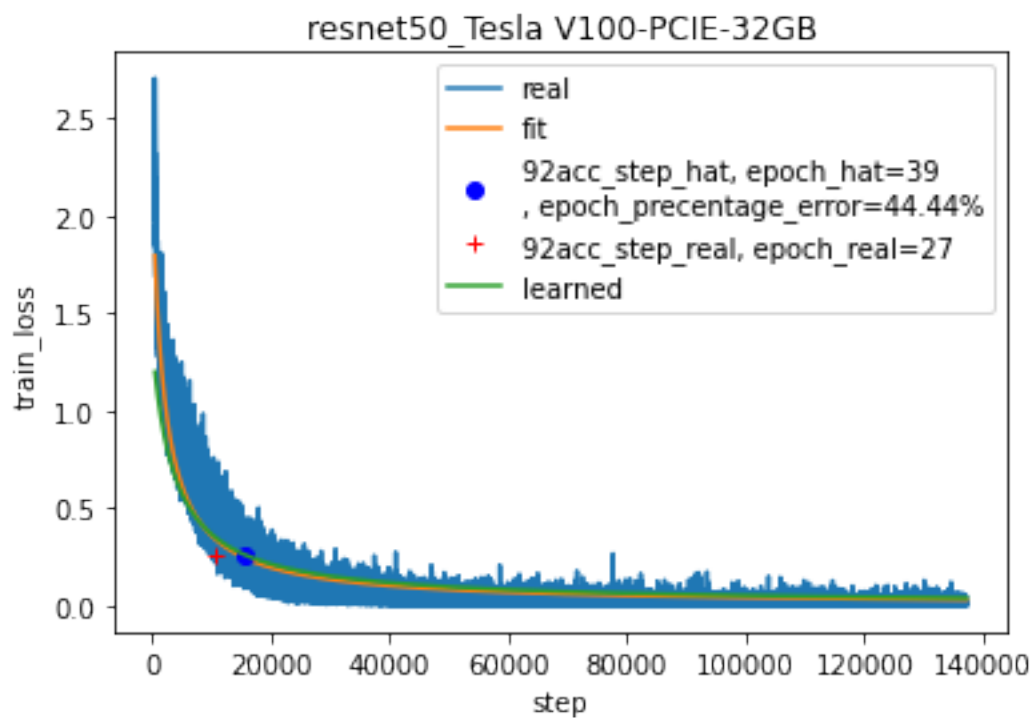
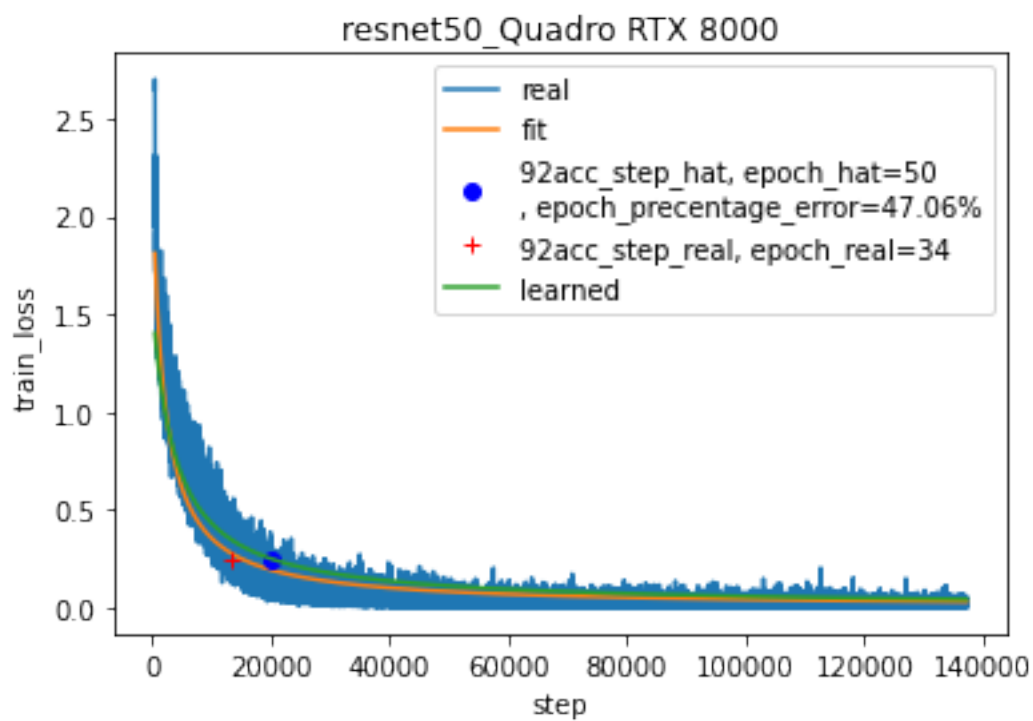
```

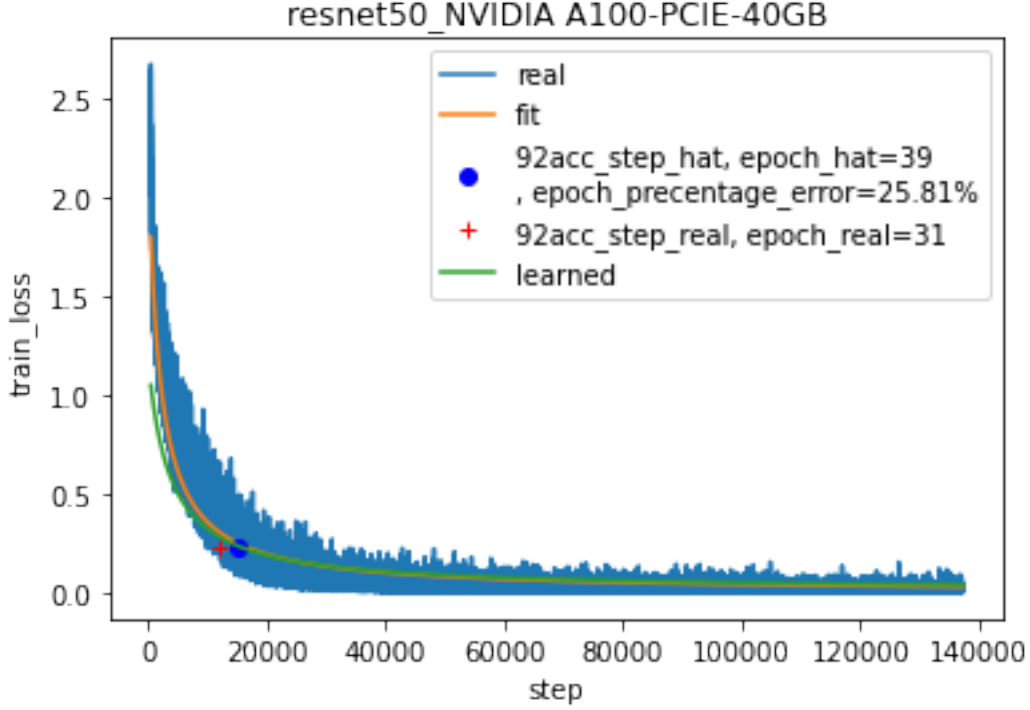
b0_hat = df["b0_hat"][i]
b1_hat = df["b1_hat"][i]
b2_hat = df["b2_hat"][i]
learned_X = real_X
learned_Y = get_y(learned_X, (b0_hat, b1_hat, b2_hat))

step_real = test_res_df["step_to_92acc"].to_list()[i]
# print(step_real)
loss_92acc_real = test_res_df["loss_92acc"].to_list()[i]
# print(step_real)
# step_hat =
threshold = test_res_df["loss_92acc"].to_list()[i]
step_hat = np.where(learned_Y==learned_Y[learned_Y<threshold][0])[0][0]
num_steps_per_epoch = 391
epoch_hat = step_hat // num_steps_per_epoch
epoch_real = step_real // num_steps_per_epoch
p = "{0:.2f}".format((epoch_hat-epoch_real)/epoch_real * 100)
plt.plot(step_hat, learned_Y[step_hat], "bo", label=f"92acc_step_hat,
↪epoch_hat={epoch_hat} \n, epoch_precentage_error={p}%")
plt.plot(step_real, loss_92acc_real, "r+", label=f"92acc_step_real,
↪epoch_real={epoch_real}")

plt.plot(learned_X, learned_Y, label="learned")
plt.xlabel("step")
plt.ylabel("train_loss")
plt.title(df["file_name"][i].replace(".csv", ""))
plt.legend()
plt.show()

```





1.3.6 Comment

The Learned curve percentage error is not that small, one reason is that we use the the kernel regression which may not extrapolate well, another reason is that the loss curve is not smooth. Actraully, the total epoch is quite small, if we check the results over all the running epoch numbers, it is quite good.

Some discussions with students: For the regression model, I use the GPU type as the categorical feature, and they may use different regression model for different GPU types, the leaned regression model maybe more accurate if we learn the model seperately.

1.3.7 3

Using the predicted number of epochs for Resnet-50 along with the resource-speed model (use Equation (4) in Peng et al. along with its coefficients from the paper) obtain the time to accuracy of Resnet-50 (to reach 92% accuracy) in two different setting (with 2 and 4 parameter servers respectively) as a function of the number of workers. So you will be plotting two curves, one for 2 and one for 4 parameter server case. Each smooth curve will show how the time to achieve 92% accuracy (on the y-axis) scales with number of workers (on the x-axis). (7)

Hint: The theta values are given in the paper. You may re-use them directly

$$f(p, w) = \left(\theta_0 \cdot \frac{M}{w} + \theta_1 + \theta_2 \cdot \frac{w}{p} + \theta_3 \cdot w + \theta_4 \cdot p \right)^{-1} \quad (1)$$

	theta_(1)	theta_(2)	theta_(3)	theta_(4)	theta_(5)	Residual sum of squares for fitting
Async	2.83	3.92	0.00	0.11	-	0.10
Sync	1.02	2.78	4.92	0.00	0.02	0.00

```
[265]: M = 128

def f_async(w, p):
    theta_0 = 2.83
    theta_1 = 3.92
    theta_2 = 0.00
    theta_3 = 0.11
    # theta_4 = 0.02
    return theta_0 * M / w + theta_1 + theta_2 * w / p + theta_3 * w

def f_sync(w, p):
    theta_0 = 1.02
    theta_1 = 2.78
    theta_2 = 4.92
    theta_3 = 0.00
    theta_4 = 0.02
    return 1 / (theta_0 * M / w + theta_1 + theta_2 * w / p + theta_3 * w +
    ↪theta_4 * p)
```

```
[266]: def get_time(w, p):
    return epoch_hat / f_sync(w, p)

def get_time_p2(w, p=2):
    return epoch_hat / f_sync(w, p)

def get_time_p4(w, p=4):
    return epoch_hat / f_sync(w, p)
```

```
[267]: w = np.arange(1, M, 1)
t2 = [get_time_p2(w_i) for w_i in w]
t4 = [get_time_p4(w_i) for w_i in w]
```

```
[268]: # plot
plt.plot(w, t2, label="p=2")
plt.plot(w, t4, label="p=4")
plt.legend()
plt.xlabel("w")
plt.ylabel("time to 92% accuracy")
plt.title("time to 92% accuracy vs number of workers(w)")
```

[268]: Text(0.5, 1.0, 'time to 92% accuracy vs number of workers(w)')

